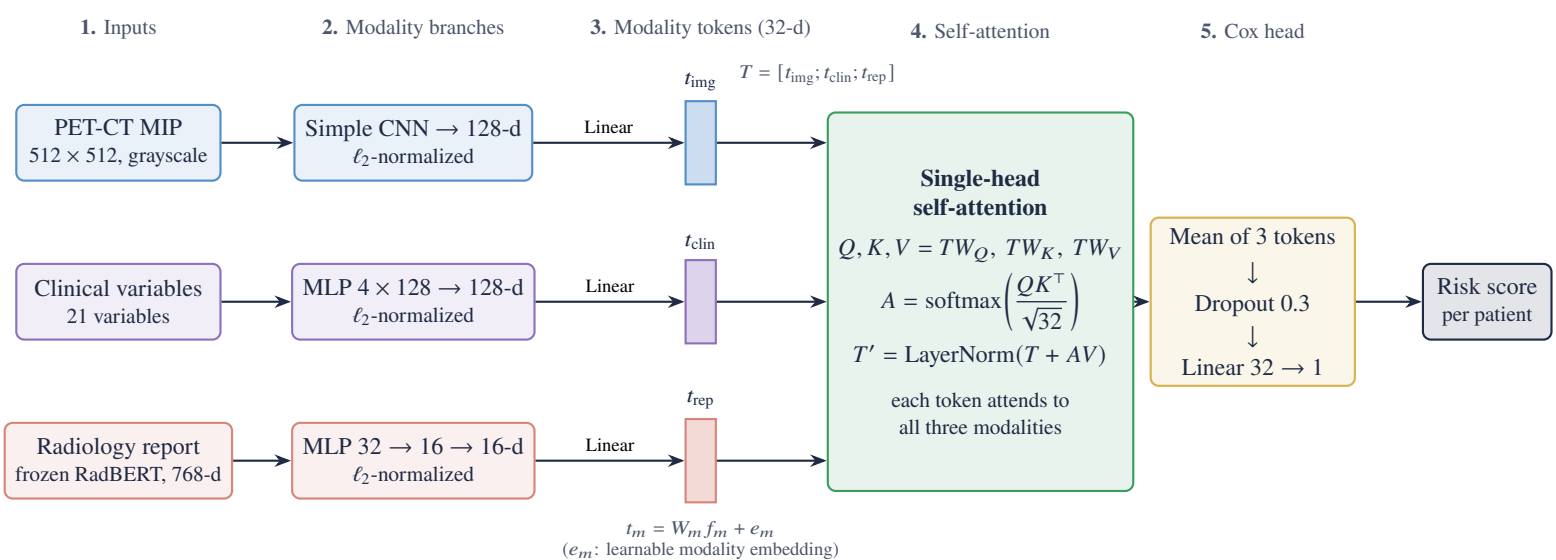


Training (end-to-end, same settings as early fusion): Cox partial likelihood; batch 32, 60 epochs, 5-fold CV (238 patients, seed 42); RadBERT frozen.
Compared with early fusion, only the mixing step differs (concatenation → self-attention); about 1.3×10^4 added parameters. C-index: OS 0.712 ± 0.037, PFS 0.648 ± 0.062.